Agentic AI Engineer

Job description

Location: Greater Vancouver Area (In-Office Only) candidates only

About D3 Security

D3 Security is transforming SecOps with Morpheus, our AI-driven Autonomous Security Operations Center (ASOC) platform. Morpheus automates Tier 1–3 analyst work with unmatched precision, processing millions of alerts in real time and empowering security teams to respond faster and smarter. Our mission: to deliver scalable, intelligent, and integrated security operations solutions for the world's most forward-thinking organizations.

The office is located in beautiful downtown Vancouver next to Burrard Skytrain Station, offering a friendly environment with an energetic, fast-growing and forward-thinking team.

Role Summary

D3 Security is looking for a hands-on Agentic AI Engineer — a software engineer who designs, builds, and manages AI agents capable of autonomous decision-making, planning, and tool use to execute complex workflows. This role sits at the intersection of agentic AI engineering, full-stack development, and AI-assisted coding workflows.

You will work on practical product problems, build reliable software, and apply LLMs and agentic patterns to real customer and internal engineering use cases. This role is ideal for someone who is strong in implementation, comfortable working across the stack, and excited to build production-grade AI agents in an enterprise environment.

Please note: this is a software engineering role, not a research or model-training position — you will build production agents on top of foundation models rather than train models. You do not need an AI/ML research background: if you are a strong software engineer who has built features with LLM APIs or used AI coding tools seriously, you are exactly who we are looking for.

What You'll Be Doing

• Design, build, and manage AI agents capable of autonomous decision-making, planning, and tool use to execute complex customer and internal workflows.
• Implement agent orchestration, multi-step planning, and tool/function calling that connect LLMs to backend services, APIs, and data systems.
• Build guardrails, validation, and human-in-the-loop controls that keep agent behavior reliable, safe, and auditable.
• Work across services and product surfaces to integrate agentic capabilities into production systems in a reliable and maintainable way.
• Collaborate with engineering, product, and design teammates to define agent behaviors, evaluate solutions, and ship user-facing improvements.
• Improve the quality of agentic systems through evaluation, testing, experimentation, and iteration — including measuring task success, latency, and cost.
• Support production readiness through sound engineering practices including observability, debugging, CI/CD, and operational reliability.
• Contribute to internal patterns and reusable components — prompts, tools, and agent scaffolding — for building agentic features effectively across the product.

What You'll Bring
• 2+ years of software engineering industry experience building and shipping production applications.
• Strong coding skills in Python and/or C#, with solid backend engineering fundamentals.
• Experience with modern web development using JavaScript/TypeScript, and ideally React or similar frameworks.
• Experience building or integrating API-driven systems and working across application layers.
• Hands-on experience building LLM-powered or agentic applications, such as tool/function calling, structured outputs, RAG, or multi-step agent workflows.
• Ability to turn ambiguous product or engineering requirements into structured implementation plans.
• Good judgment around when agents should act autonomously, and when deterministic software logic, validation, or human approval should drive behavior.
• Understanding of core agent design patterns: planning, memory and state management, tool use, and error handling and recovery.
• Attention to quality, a builder mindset, and comfort with code reviews, testing, and continuous improvement
• Clear written and verbal communication.

Nice to Have

• Has used AI coding tools to meaningfully change how they build software, not just experimented with them.
• Familiarity with Docker, Kubernetes, or Azure
• Experience shipping AI-powered product features in enterprise software
• Interest in cybersecurity, automation, or workflow-heavy systems
• Exposure to evaluation, prompt iteration, or reliability practices for LLM applications

Compensation & Benefits

• We offer competitive compensation along with performance-based bonuses and equity opportunities for exceptional performers.
• Extended health and dental benefits
• On-site gym access
• Commuter expense reimbursement
• Team events and social activities
• Free foods and snacks in-office
• Most importantly: