AI Evaluation Engineer (Remote)

About the Project Outlier helps the world’s most innovative companies improve their AI agents by providing human feedback. Do you want to shape the future of autonomous agents like OpenClaw? We collaborate with leading AI organizations to train Large Language Models (LLMs) to function as proactive, multi-step agents.

Our projects focus on teaching these systems how to design, coordinate, and optimize complex, real-world architectural workflows. Whether you are a passionate orchestration guru or experienced software developer — we want you to help us train the world's most advanced generative systems.

Ideal

Qualifications • 2+ years of experience in backend engineering, AI automation, or complex systems integration.
• Proven ability to build and maintain production-grade software with modular separation (e.g., distinct services for data parsing, logic processing, and reporting).
• Strong command of at least two major languages (e.g., Python, JavaScript, Go, or Java) and experience working with SQL databases.
• Practical experience building for live, non-mocked environments and handling multi-turn system interactions.
• Outstanding attention to detail and the ability to provide clear, high-density technical feedback on complex system behaviors. Nice to have • Expertise building multi-stage coordination tasks where data acquisition leads to reasoned output.
• Hands-on experience integrating agents with live tools such as Supabase, Gmail, and various APIs to solve real-world problems.
• High level of comfort implementing persistent state and session discovery using MEMORY.md to track agent progress.
• Experience identifying subtle failures like privacy leaks, authority escalation, or indirect prompt injections.