AI Adversarial Specialist - Fully Remote

Mercor

About the job

About The Job

Mercor connects elite creative and technical talent with leading AI research labs. Headquartered in San Francisco, our investors include Benchmark, General Catalyst, Peter Thiel, Adam D'Angelo, Larry Summers, and Jack Dorsey.

Position: AI Safety Experts — English & Bengali

Type: Contract

Compensation: $20–$22/hour

Location: Remote

Role Responsibilities

- Red team conversational AI models and agents. Focus on jailbreaks, prompt injections, misuse cases, and bias exploitation.
- Generate high-quality human data. Annotate failures, classify vulnerabilities, and flag systemic risks.
- Apply structure using taxonomies, benchmarks, and playbooks for consistent testing.
- Document reproducibly. Produce reports, datasets, and attack cases for customer action.
- Work independently and asynchronously. Ensure flexibility and adaptability across projects.

Qualifications

Must-Have

- Fluent in English and Bengali.
- Prior red teaming experience in AI adversarial work, cybersecurity, or socio-technical probing.
- Strong communication skills for explaining risks to technical and non-technical stakeholders.

Preferred

- Experience with adversarial ML, cybersecurity, and socio-technical risk.
- Skills in creative probing like psychology, acting, or writing for unconventional adversarial thinking.

Application Process (Takes 20–30 mins to complete)

- Upload resume
- AI interview based on your resume
- Submit form

Resources & Support

- For details about the interview process and platform information, please check: https://talent.docs.mercor.com/welcome
- For any help or support, reach out to: support@mercor.com

PS: Our team reviews applications daily. Please complete your AI interview and application steps to be considered for this opportunity.

,