About the Role

The Grokipedia team at xAI is building an AI-powered knowledge platform that leverages advanced LLMs to generate, maintain, and update high-quality, factual articles. We are seeking a skilled engineer to lead the development and improvement of article generation pipelines, focusing on epistemics to ensure truthfulness and reliability. This role involves creating LLM agents for long-form content creation, synthetic data generation for training and evaluation, advanced prompting techniques, and inference systems that incorporate user feedback and keep content current.

Tech Stack

- Python
- Agent frameworks
- Ray

Focus

- Developing article rewrite pipelines using LLM agents to produce accurate, long-form content while prioritizing epistemic integrity.
- Developing prompting strategies and synthetic data generation methods to improve model performance in knowledge-intensive tasks.
- Incorporating user edit suggestions into inference pipelines to iteratively refine articles for factual accuracy and completeness.
- Building systems to automatically update articles with new information, ensuring ongoing truthfulness and reducing hallucinations through rigorous verification mechanisms.
- Experimenting with epistemics-focused techniques, such as chain-of-thought reasoning, self-critique, and multi-agent debates, to elevate the quality of generated content.
- Collaborating on scalable infrastructure for real-time article maintenance and generation at production scale.

Ideal Experiences (at least one from below)

- Strong engineering background with expertise in LLMs, agentic systems, and epistemics in AI.
- Experience in prompting engineering, synthetic data creation, and long-form text generation using models
- Proficiency in building and optimizing inference pipelines for knowledge bases, including handling user feedback loops.
- Familiarity with techniques for ensuring factual accuracy in AI outputs, such as retrieval-augmented generation (RAG) or epistemic uncertainty modeling.
- Hands-on work with large-scale data processing and generation, focusing on quality, verifiability, and bias mitigation.
- Preference for candidates with experience in content platforms, wiki-like systems, or AI-driven knowledge management tools.

Location

The role is based in the Bay Area [San Francisco and Palo Alto]. Candidates are expected to be located near the Bay Area or open to relocation.

Interview Process

After submitting your application, the team reviews your CV and statement of exceptional work. If your application passes this stage, you will be invited to a 15-minute interview (“phone interview”) during which a member of our team will ask some basic questions. If you clear the initial phone interview, you will enter the main process, which consists of four technical interviews:

- Coding assessment in a language of your choice.
- Systems hands-on: Demonstrate practical skills in a live problem-solving session.
- Project deep-dive: Present your past exceptional work to a small audience.
- Meet and greet with the wider team.

Our goal is to finish the main process within one week. All interviews will be conducted via Google Meet.

Annual Salary Range

$180,000 - $440,000 USD