[Hiring] [Remote] [US] - Senior Machine Learning Engineer, AI Platform ($139k-$218k)

HIRING

The AI Platform team is responsible for building the foundational infrastructure that powers intelligent experiences across Mozilla products. This includes model training pipelines, high-throughput inference services, GPU orchestration, and secure, privacy-respecting AI systems that operate reliably at global scale.

- We're looking for a Machine Learning Engineer with a strong platform mindset to help design, build, and operate Mozilla's AI platform. In this role, you'll work at the intersection of machine learning, distributed systems, and production infrastructure—ensuring that models can be trained, deployed, and served efficiently, securely, and at scale. You will collaborate closely with product, infrastructure, and security teams to enable fast iteration while meeting strict performance and privacy requirements.

What You'll Bring:

- Bachelor's degree with 4–6 years of relevant industry experience, or Master's degree with significant hands-on experience building and operating production ML systems, or work experience equivalent
- Strong experience developing in Python for machine learning systems, backend services, or distributed data processing.
- Proven experience deploying and operating ML workloads in cloud environments, including production-grade infrastructure.
- Solid understanding of model serving architectures, inference pipelines, and performance tradeoffs (latency, throughput, cost, scaling strategies).
- Hands-on experience working with GPU-based workloads and accelerated computing in production settings.
- Experience designing CI/CD pipelines and development workflows that support reliable ML system deployment.
- Ability to independently scope and drive technical initiatives while balancing product and operational priorities.
- Strong problem-solving skills and the ability to debug performance and reliability issues in distributed systems.
- Clear and effective communication skills, with experience collaborating across engineering, product, and infrastructure teams.

Bonus Skills:

- Experience implementing inference optimization strategies such as batching, quantization, compilation, model conversion, or hardware-specific tuning.
- Familiarity with containerization and orchestration systems (e.g., Docker, Kubernetes) in production environments.

- Experience designing observability systems for distributed services, including metrics strategy and performance profiling.
- Exposure to privacy-preserving ML techniques, security best practices, or responsible AI system design.
- Contributions to open-source ML infrastructure projects or leadership in building reusable internal ML tooling.