Expert Professionals — AI/ML

Location Requirements

United States
Location: Remote
Description

Join a leading AI lab's cutting-edge GenAI team to be at the core of the AI revolution, where your expertise fuels the development of the most advanced Large Language Models.

1. Overview

A leading AI lab is advancing the frontier of agentic AI systems by building models that can reason through, execute, and evaluate complex technical tasks across domains like machine learning, data science, data engineering, and software development. We're seeking talented AI & Data Science subject-matter experts (SMEs) to bring deep technical expertise and elevate the quality of our AI training data.

This is a W-2 employment position with Cincinnatus LLC, requiring a commitment of 40 hours per week (during weekdays). This position will be placed at a leading AI Lab as part of their extended workforce.

2. Key Responsibilities

- Guide research and engineering teams to close knowledge gaps in AI and data science domains, surfacing the nuances that distinguish expert-level work from surface-level reasoning.
- Design challenging agentic tasks rooted in real-world ML, data science, data engineering, and software workflows, and write accurate, well-documented solutions that serve as ground truth.
- Evaluate AI agent outputs against your solutions, providing detailed written feedback that captures correctness, efficiency, and reasoning quality.
- Develop and refine evaluation frameworks and rubrics for assessing agentic behavior on AI and data science tasks.
- Collaborate with other subject matter experts to ensure consistency and accuracy in training data.

3. Core Qualifications

- 3+ years of research, academic, or industry experience in Machine Learning, or Data Science, or Software Engineering, or Computer Science, or Statistics, or Biology, or Electrical/Mechanical/Civil Engineering, or Physics, or Chemistry, or Mathematics, or Materials Science, or other STEM background.
- Demonstrated technical expertise in at least one of: programming, data analysis, ML modeling, statistical methods, or computational methods.
- Ability to commit to 40 hours per week during weekdays for the duration of the engagement.
- Prior experience with data annotation, labeling, evaluation, or human feedback collection is a strong plus.
- Experience with LLMs, AI systems, or agentic workflows; familiarity with agentic frameworks is a plus.
- Strong written communication skills and the ability to explain technical decisions clearly.