Detection & Response Engineer
Palo Alto, CA

ABOUT THE ROLE:
xAI was first in the world to build a 100k GPU cluster on an ethernet network and then did it again in 92 days, floors, walls and all. We need an engineer with deep experience in RoCEv2 that can develop at hyper scale while optimizing performance and availability.

xAI is building at a furious pace with the latest hardware to help people understand the universe. To make the next significant leap forward, we need to own our own destiny by understanding our current network performance and availability and then optimize it to our training models and how we execute customer inference queries. You will spend most of your days deep inside NCCL, building metric dashboards and tweaking configurations to ensure no performance is left on the table. You will help design the next iteration of our backend and front-end networks that will allow us to seamlessly build-out new GPU infrastructure with little to no engineering assistance.

There will be a significant amount of travel to Memphis for building more capacity as well as participating in a team on-call rotation and helping on other scaling and maintenance efforts. This will become easier as we build out the team and engineers contribute to deployment and operations frameworks to remove repetitive tasks.

BASIC QUALIFICATIONS:
A minimum of 10 years designing and operating large scale networks with 5 years in the ethernet AI/HPC space.
Deep understanding of congestion control on ethernet with Infiniband an added bonus.
Deep understanding of AI training and inference workloads and how they operate on the network. As part of this you are able to use and debug NCCL and potentially commit to the library.
Expertise in creating a portfolio of metrics for performance and operations to optimize the fleet for training and inference traffic.
Experience with Python to automate away repetitive tasks and facilitate your daily job working with and analyzing large sets of data.